



Meet Kafka

“the distributed log”



BWM- 27 octobre 2017 - Lyon



Florian Garcia

Software engineer at Ippon



fgarcia@ippon.fr



[@lm_flog](https://twitter.com/lm_flog)

- **Backend lover**
- **Zipkin committer in spare time**
- **Interested in**
 - Distributed tracing
 - Kafka
 - Docker
- **Speaker**



Kafka

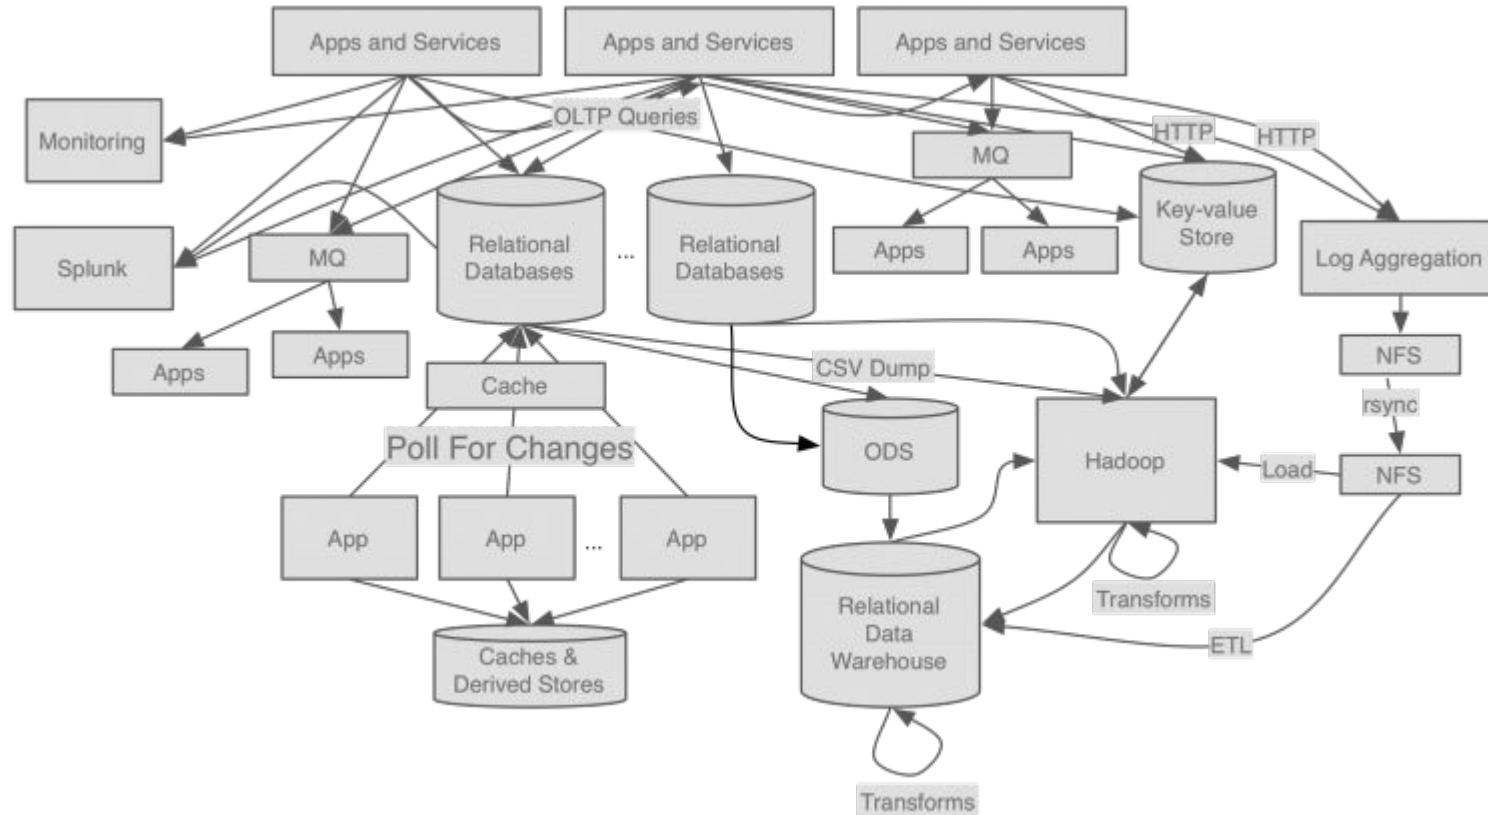
- Kafka who are you?
- Kafka use cases
- Kafka 101
- Kafka Streams
- Kafka from the trenches



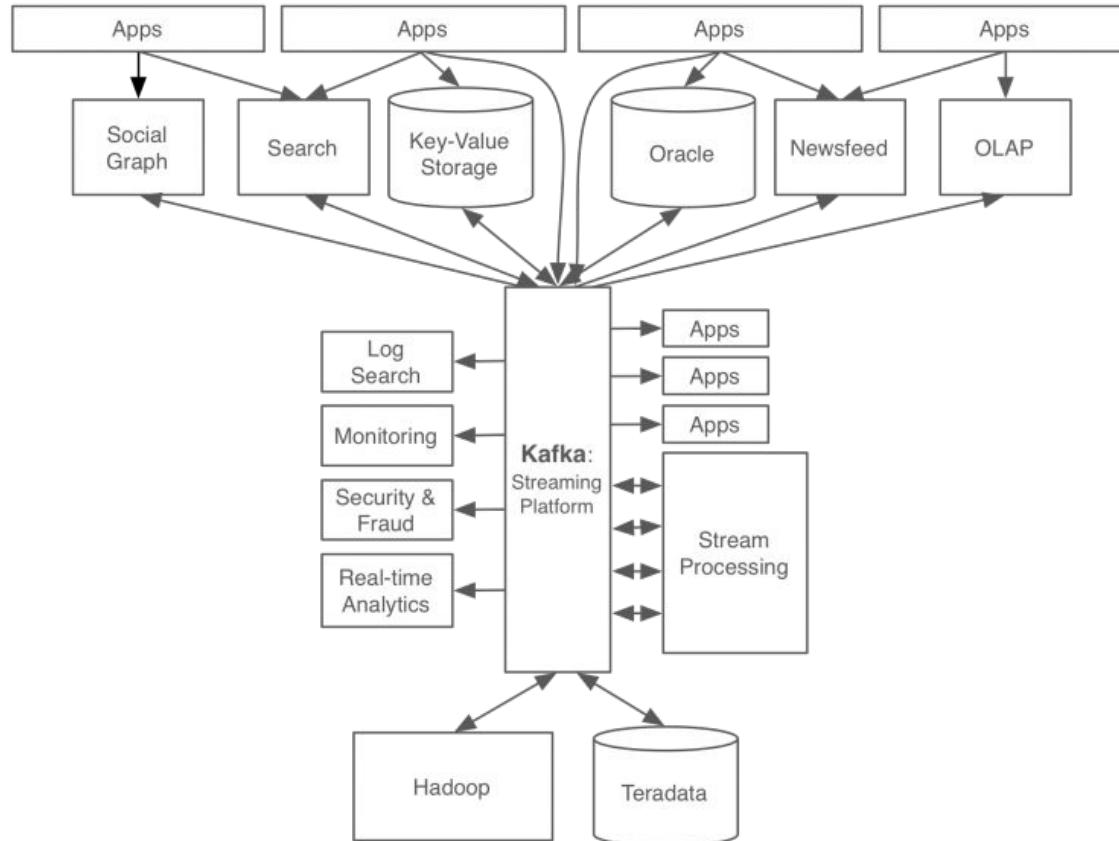
Kafka, who are you?



Linkedin before Kafka (pre 2010)

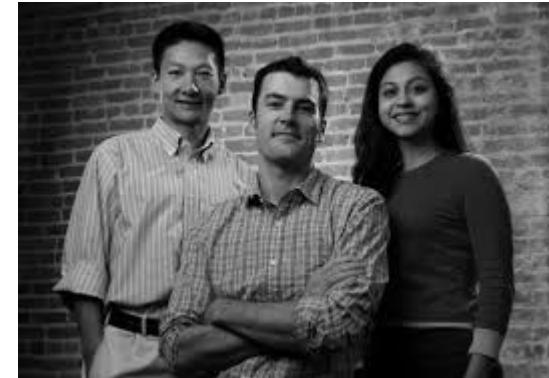


Linkedin after Kafka



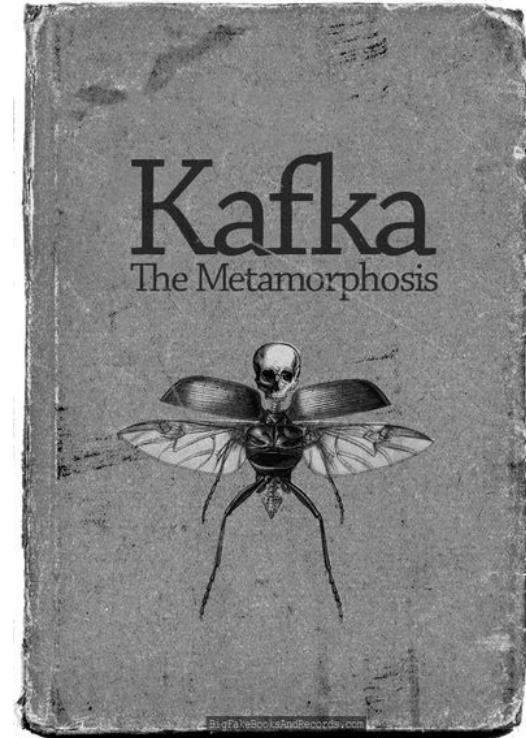
What about today ?

- Open sourced in 2011
- Now a top level Apache project
 - Very active
 - Current release 0.11.0.X
- Spinoff company Confluent created in 2014
 - Jay Kreps, Neda Narkhede and Jun Rao
 - Ships the Confluent platform
 - Recently raised \$50M in a round C funding



The metamorphosis

- Building a Kafka ecosystem
 - Kafka Connect
 - Schema Registry
 - Integrations
- A lot of new features regularly added
- From distributed log to a streaming platform
- End to end exactly once processing !



Glossary

- Broker \Leftrightarrow a server running Kafka
- Topic \Leftrightarrow category where records are published
- Partition \Leftrightarrow shard of a topic
- Producer \Leftrightarrow client writing to Kafka
- Consumer \Leftrightarrow client reading from Kafka





Kafka Use Cases

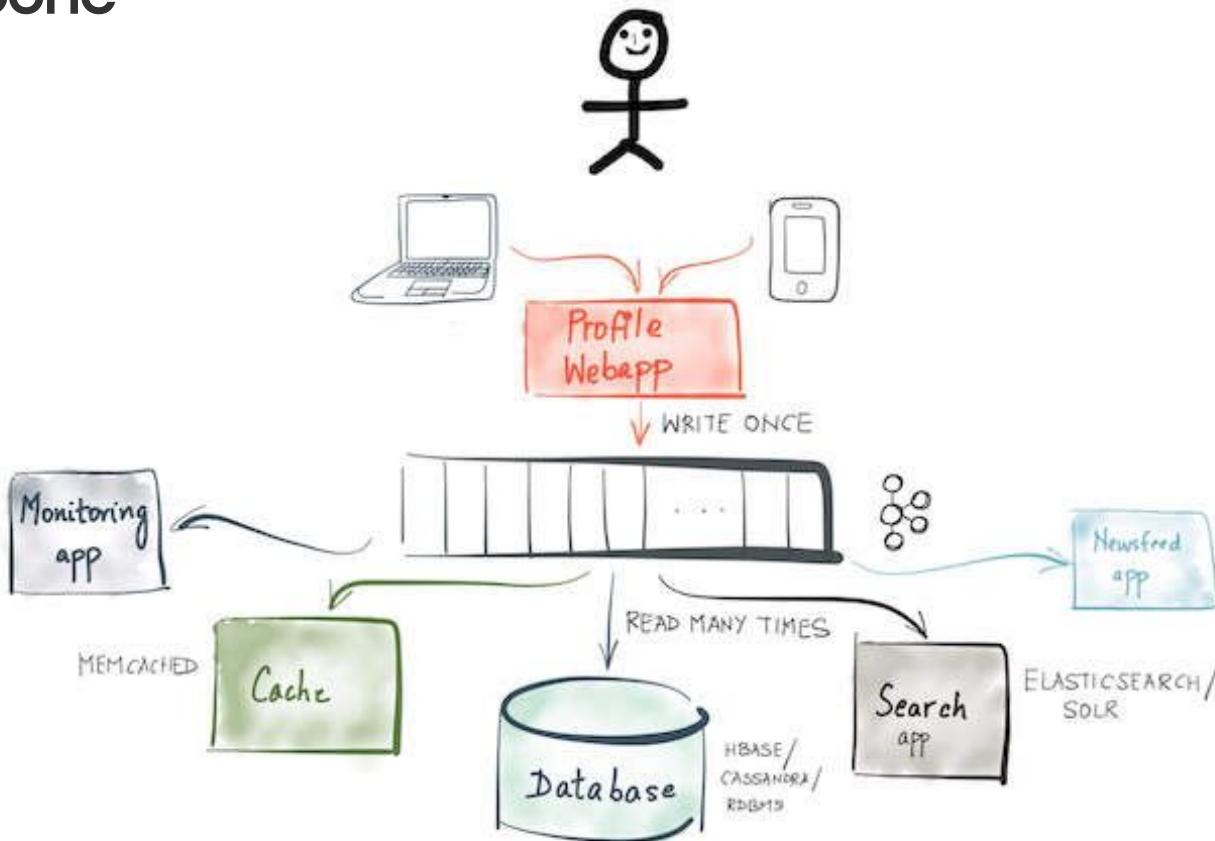
PPON

Data deluge

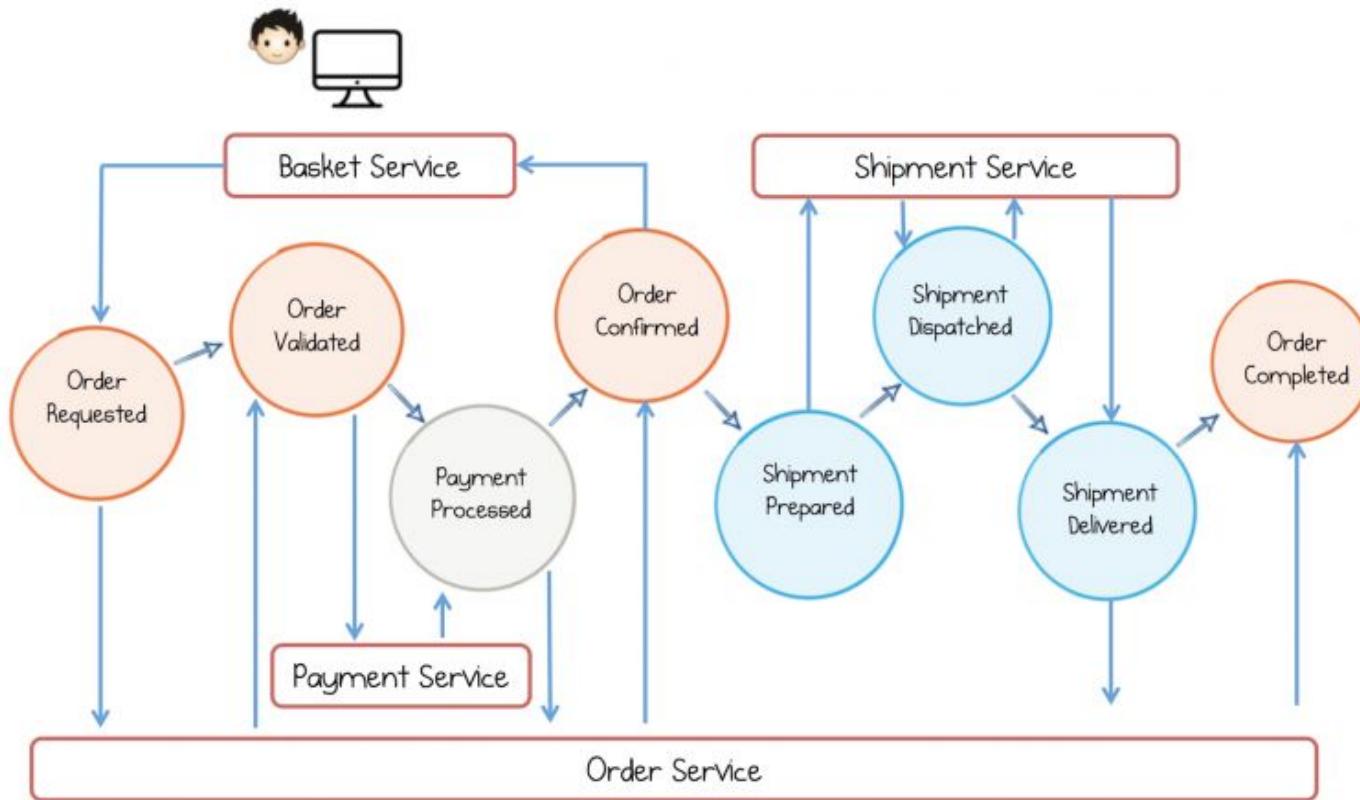
- User activity
- Server info
- Messaging
- Business data consolidation
- Event sourcing



As a backbone

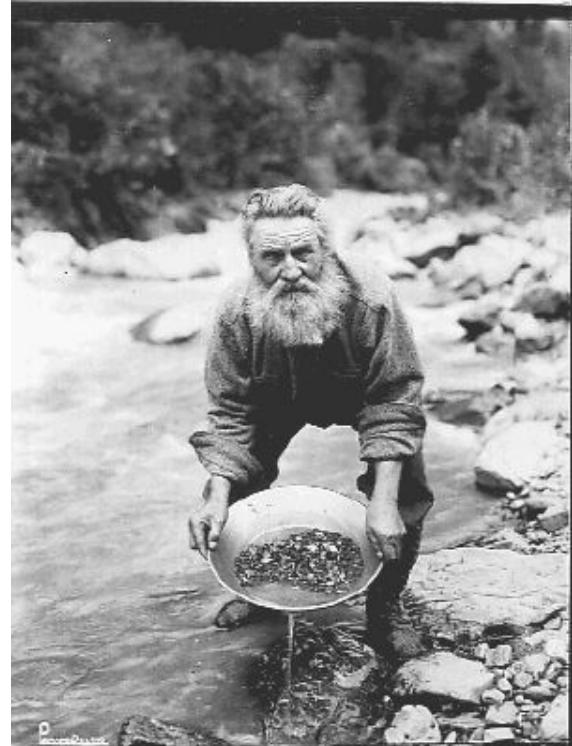


For event driven architectures



Streaming !

- Process data as soon as they are available !
- Powerful
- Easy to use
- Real time
 - decision
 - data consolidation
 - monitoring
 - fraud detection
 - ...



But wait, I already bought an ESB

- ESB
 - complex logic
 - routing
 - processing
- Smart endpoints, dumb pipes !
- Kafka just transports the data
- Get rid of the ESB before you need to sell your children



But wait, I already use RabbitMQ

- **RabbitMQ**
 - Broker centric
 - Complex routing
 - Mature (Erlang)
 - Standardized (AMQP)
- **Kafka**
 - Consumer / producer centric
 - Ordered (per partition)
 - Higher performance
 - Backed by confluent



Who uses Kafka ?



ORACLE



Goldman
Sachs



<https://kafka.apache.org/powerd-by>

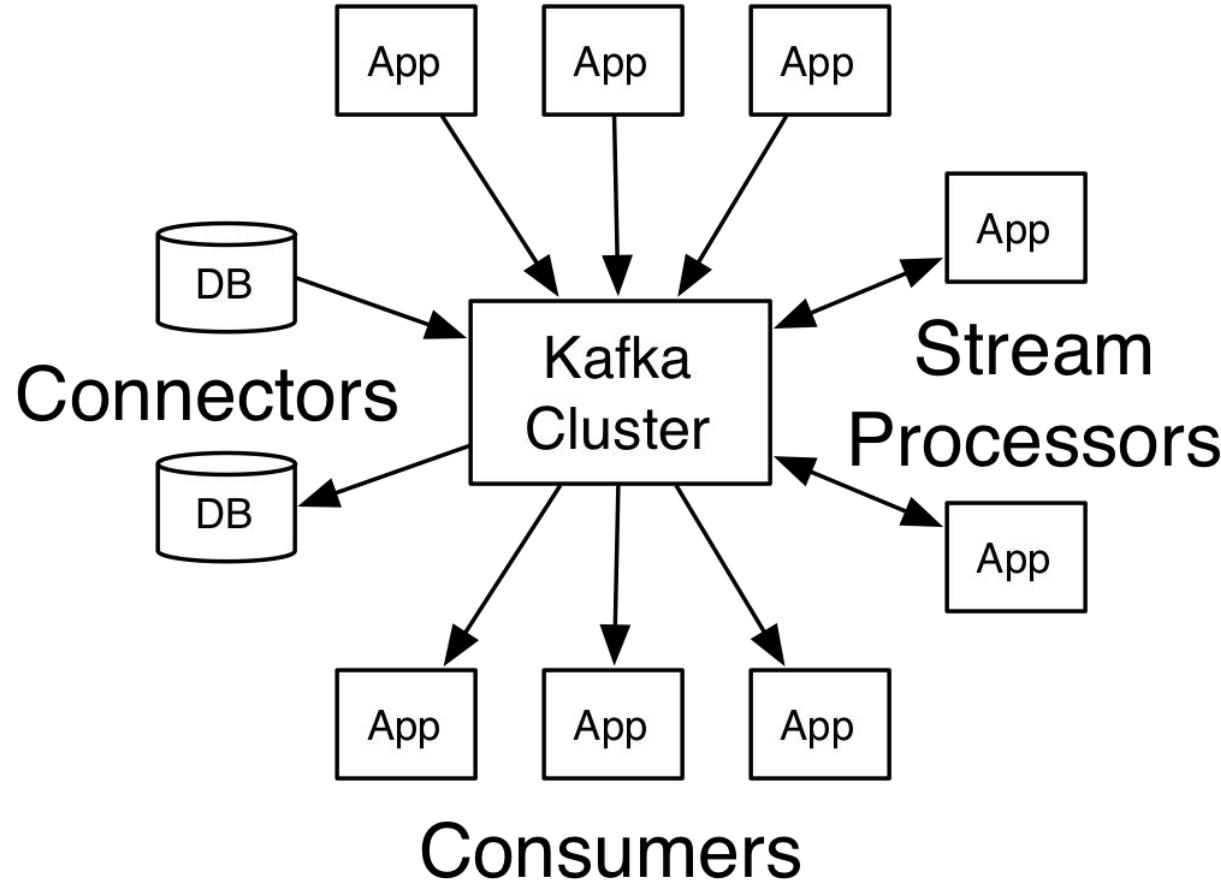


Kafka 101

PPON

Kafka centric

Producers



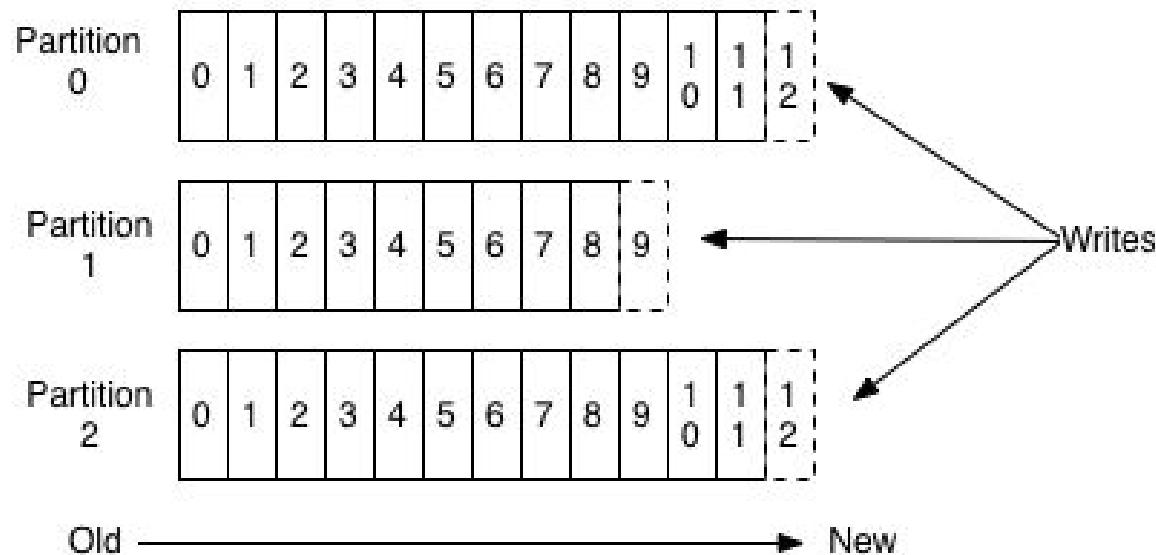
Brokers

- A running Kafka service
- Role
 - Read / write on disk
 - Consumer / producer interface
 - Group balancing, leader election
- Uses Zookeeper for coordination



Topics and partitions

Anatomy of a Topic



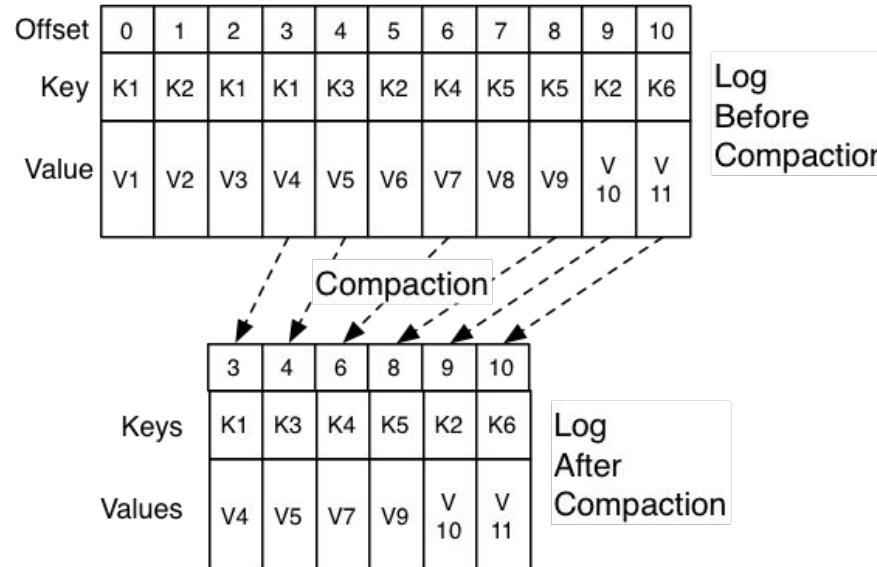
Message content

- **RecordBatch**
 - FirstOffset & LastOffsetDelta
 - FirstTimestamp & MaxTimestamp
 - MagicByte => 2 for version > 0.11.X
 - Records
- **Record**
 - Timestamp & Offset deltas
 - Key
 - Value
 - Headers



Compaction

- Keep only the last log for a given key
- Happens after a new segment creation by the log cleaner



Kafka guarantees

Understanding Streaming Semantics



At most once	At least once	Exactly once
Message pulled once	Message pulled one or more times; processed each time	Message pulled one or more times; processed once
May or may not be received	Receipt guaranteed	Receipt guaranteed
No duplicates	Likely duplicates	No duplicates
Possible missing data	No missing data	No missing data

Kafka Clients

- Low level libraries
 - Java
 - C/C++ (librdkafka)
- Built on top of librdkafka
 - C#
 - Go
 - Haskell
 - Python
 - ...

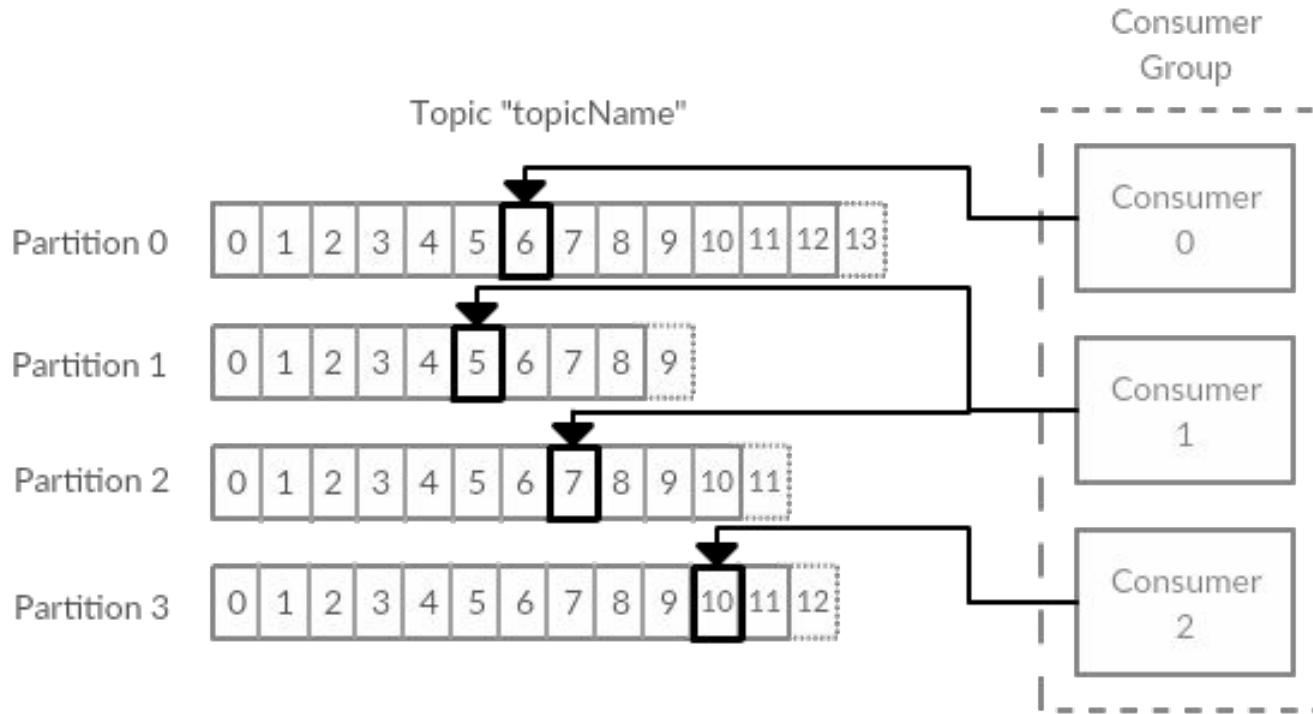


Producing

- Send data to topics
- Binary content
- The key define the partition
- Send asynchronously
 - buffer pending records
 - batch records for efficiency
- Retry in case of failure



Consuming



Dataset - french opendata



ENSEIGNEMENT SUPÉRIEUR
ET RECHERCHE

[Accueil](#) | [Explorez les jeux de données](#) | [Notre démarche](#) | [Réutilisations](#) | [Créez vos cartes](#) | [Utilisez les API](#)

274 690
enregistrements

Aucun filtre actif.

Filtres

Rentrée universitaire

2015-16	19 472
2014-15	19 281
2013-14	19 118
2012-13	18 371
2011-12	18 329
2010-11	18 276
> Plus	

Niveau géographique

Commune	136 026
Unité urbaine	75 559
Département	37 651
Académie	15 646
Région	9 031
Pays	777

Effectifs d'étudiants inscrits dans les établissements et les formations de l'enseignement supérieur

Informations

Tableau

Export

API

Ce jeu de données est sous licence : Licence Ouverte (Etalab)

Formats de fichiers plats

CSV

Jeu de données entier

Le CSV utilise le point-virgule (;) comme séparateur.

JSON

Jeu de données entier

Excel

Jeu de données entier

Example : producer (settings)

```
private static Properties buildKafkaProps(boolean withTransaction) {  
    Properties props = new Properties();  
    props.put("bootstrap.servers", "localhost:9092");  
    props.put("key.serializer", "org.apache.kafka.common.serialization.IntegerSerializer");  
    props.put("value.serializer", "org.apache.kafka.common.serialization.StringSerializer");  
  
    // Transaction support  
    if (withTransaction) {  
        props.put("transactional.id", "myAppId");  
    }  
    // Exactly once  
    props.put("enable.idempotence", true);  
    return props;  
}
```

<https://github.com/lmFlog/kafka-playground>

Example : producer (without transactions)

```
try (KafkaProducer<Integer, String> kafkaProducer = new KafkaProducer<>(props)) {  
    kafkaProducer.send(  
        new ProducerRecord<>(  
            EFFECTIFS_TOPIC,  
            effectif.getYear(), // Shard by year  
            jsonMapper.writeValueAsString(effectif)));  
}
```

<https://github.com/lmFlog/kafka-playground>

Example producer (with transactions)

```
try (KafkaProducer<Integer, String> kafkaProducer = new KafkaProducer<>(props)) {  
    kafkaProducer.initTransactions();  
  
    kafkaProducer.beginTransaction();  
    kafkaProducer.send(  
        new ProducerRecord<>(  
            EFFECTIFS_TOPIC,  
            effectif.getYear(), // Shard by year  
            jsonMapper.writeValueAsString(effectif)));  
    kafkaProducer.commitTransaction();  
}
```

<https://github.com/lmFlog/kafka-playground>

Exemple : consumer

```
private static Properties buildKafkaProps() {
    Properties props = new Properties();
    props.put("bootstrap.servers", "localhost:9092");
    props.put("key.deserializer", "org.apache.kafka.common.serialization.IntegerDeserializer");
    props.put("value.deserializer", "org.apache.kafka.common.serialization.StringDeserializer");
    // Consumer group ID
    props.put("group.id", "basic");

    // Enabling transaction support
    props.put("isolation.level", "read_committed");
    return props;
}
```

<https://github.com/lmFlog/kafka-playground>

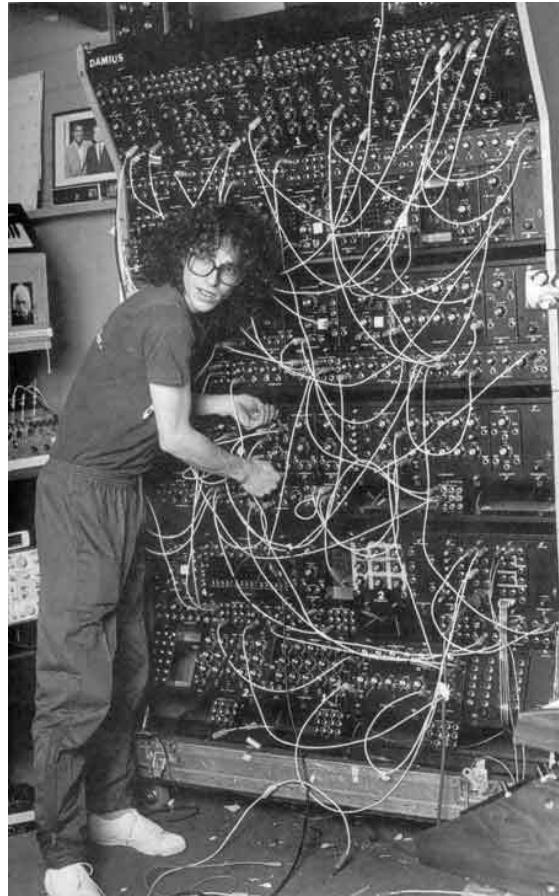
Exemple : consumer

```
KafkaConsumer<String, String> consumer = new KafkaConsumer<>(props);
consumer.subscribe(Collections.singleton(EFFECTIFS_TOPIC));
while (running) {
    ConsumerRecords<String, String> records = consumer.poll( timeout: 100 );
    for (ConsumerRecord<String, String> record : records) {
        Effectif effectif = jsonMapper.readValue(record.value(), Effectif.class);
    }
}
```

<https://github.com/lmFlog/kafka-playground>

Kafka Connect

- Standardizes integration with other systems
 - Community connectors
 - Custom connectors
- Working modes
 - Distributed
 - Standalone
- Automatic offset management
- Distributed and scalable by default

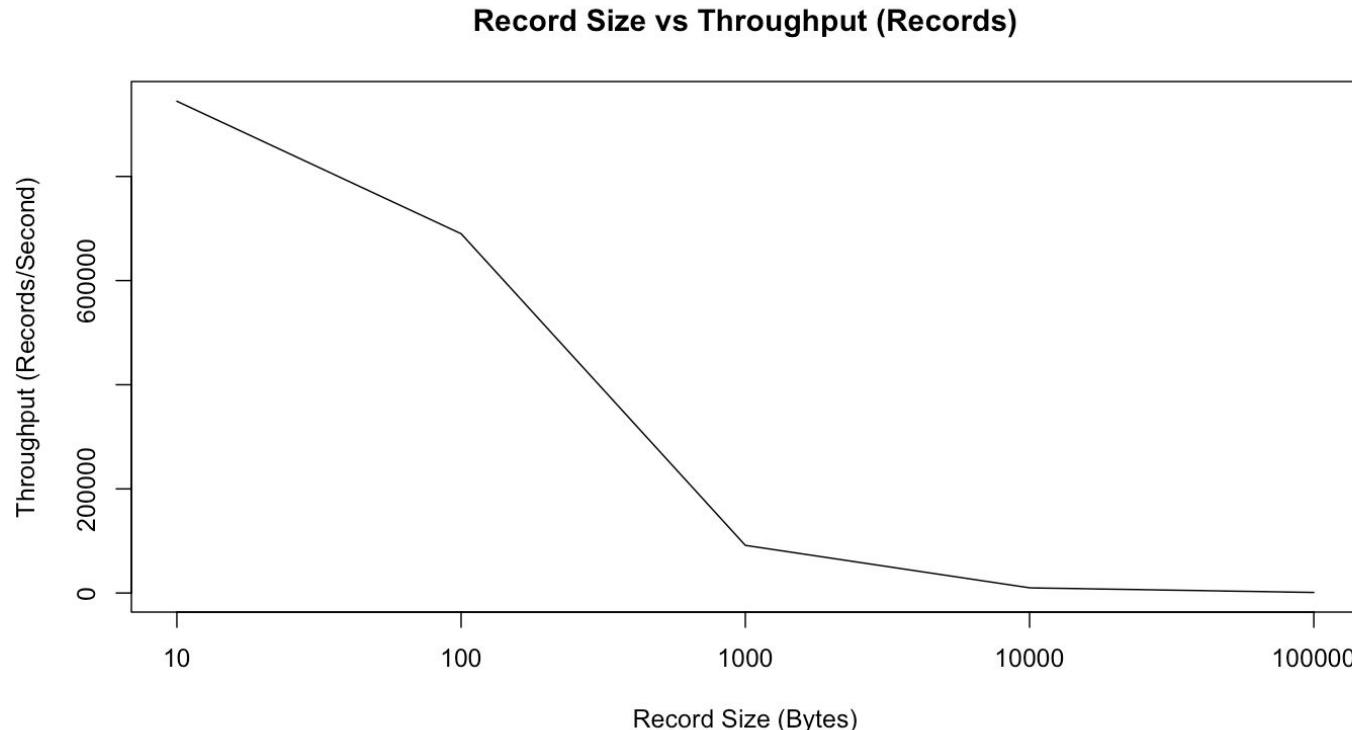


Confluent platform additions

- Schema Registry
- REST Proxy
- Connector REST interface (\$)
- Additional connectors (\$)
- Kafka replicator (\$)
- Support (\$)



Performance



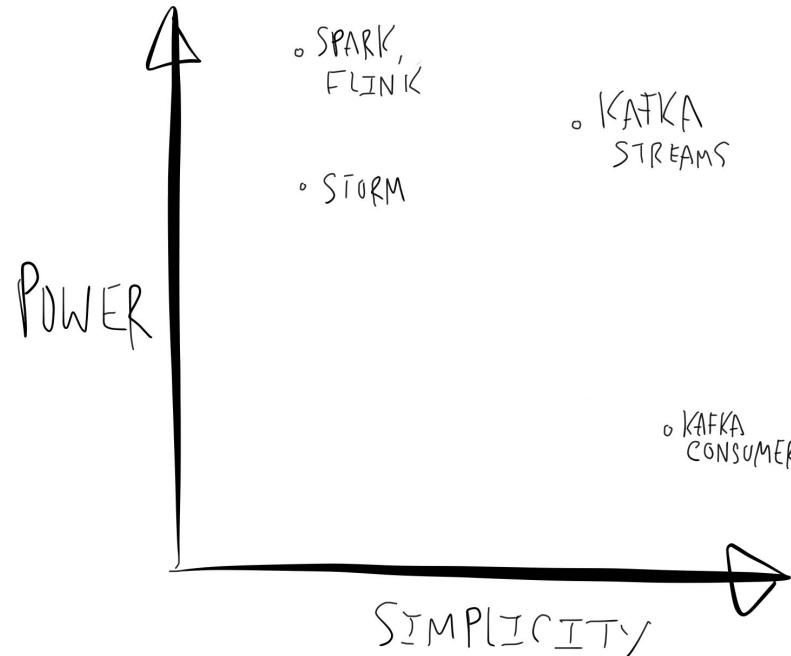


Kafka Streams



MakeAGIF.com

The goal



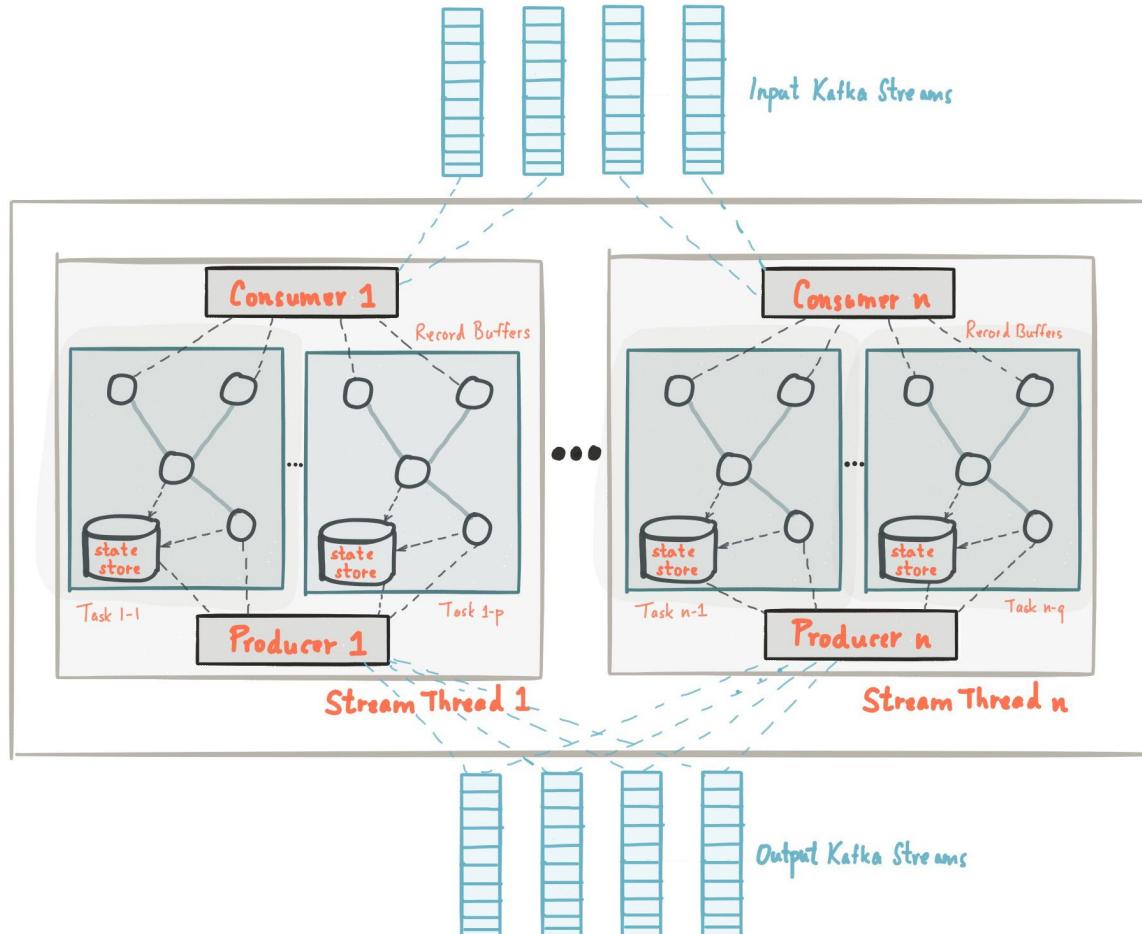
Source www.confluent.io

Kafka streams

- Available since version 0.10.0
- A simple and lightweight client library
- Event-at-a-time processing
- Millisecond latency
- Stateful processing
- Exactly once processing !



Anatomy



Source: Kafka documentation

KStream

- An abstraction of a record stream
 - Each data is an “insert”
 - No record replaces an existing row with the same key
- Example
 1. ("Blend", 1)
 2. ("Blend", 3)
 3. Sum("Blend") = ?
 4. Sum("Blend") = 4



KTable

- An abstraction of a changelog stream
 - Each data is an “upsert”
 - Based on the record key
- Example
 1. ("Blend", 1)
 2. ("Blend", 3)
 3. $\text{Sum}(\text{"Blend"}) = ?$
 4. $\text{Sum}(\text{"Blend"}) = 3$



For developers

- Java only
- kafka-stream dependency
- A high level API
 - Use java 8 lambdas
 - Map, filter, join, ...
- A low level API
 - Add and connect processors
 - Interact directly with state stores



Operations

- **Stateless transformations**
 - map
 - groupBy
 - filter
- **Stateful transformations**
 - Windowing
 - Joins
 - Aggregations



Example

```
KTable<Integer, Integer> studentsPerYear = kStreamBuilder
    .stream(Serdes.Integer(), effectifSerde, inputTopic)
    .filter((year, effectif) -> !"TOTAL".equals(effectif.getGroup()))
    .filter((year, effectif) -> "Commune".equals(effectif.getGeographicLevel()))
    .map((year, effectif) -> new KeyValue<>(year, effectif.getStudentCount()))
    .groupByKey(Serdes.Integer(), Serdes.Integer())
    .aggregate(
        () -> 0,
        (aggKey, value, aggregate) -> aggregate + value,
        Serdes.Integer(),
        s: "studentsPerYear");
    studentsPerYear.to(Serdes.Integer(), Serdes.Integer(), s: "studentsPerYearTopic");
```

Example

```
KTable<String, Integer> studentsPerCommune2014 = new KStreamBuilder()
    .stream(Serdes.Integer(), effectifSerde, inputTopic)
    .filter((year, effectif) -> !"TOTAL".equals(effectif.getGroup()))
    .filter((year, effectif) -> "Commune".equals(effectif.getGeographicLevel()))
    .filter((year, effectif) -> year == 2014)
    .groupBy((year, effectif) -> effectif.getGeographicUnit(), Serdes.String(), effectifSerde)
    .aggregate(
        () -> 0,
        (aggKey, value, aggregate) ->
            aggregate + (value.getStudentCount() != null ? value.getStudentCount() : 0),
        Serdes.Integer(),
        s: "studentsPerCommune2014");
```

```
studentsPerCommune2014
    .leftJoin(studentsPerCommune2015, this::calculateEvolution)
    .toStream()
    .print(Serdes.String(), Serdes.String(), s: "Evolution 2014 2015");
```

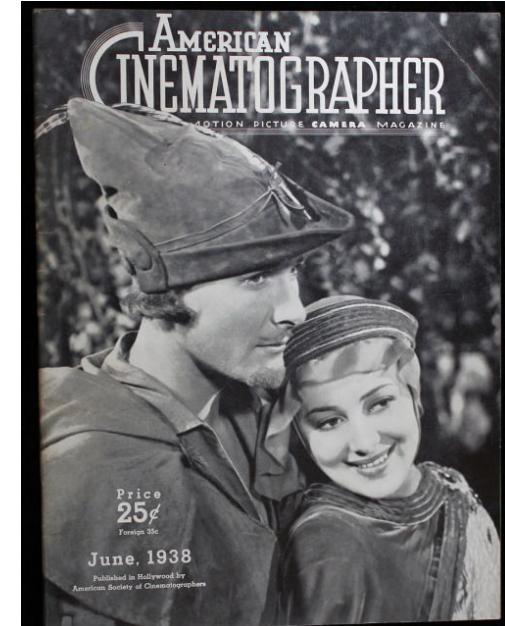
Start streaming

```
private static Properties getProperties() {  
    Properties settings = new Properties();  
    settings.put(StreamsConfig.APPLICATION_ID_CONFIG, STREAM_APPLICATION_NAME);  
    settings.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");  
    settings.put(StreamsConfig.DEFAULT_KEY_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());  
    settings.put(StreamsConfig.DEFAULT_VALUE_SERDE_CLASS_CONFIG, Serdes.String().getClass().getName());  
  
    settings.put(StreamsConfig.PROCESSING_GUARANTEE_CONFIG, StreamsConfig.EXACTLY_ONCE);  
    return settings;  
}
```

```
StreamsConfig config = new StreamsConfig(getProperties());  
KafkaStreams kafkaStreams = new KafkaStreams(kStreamBuilder, config);  
kafkaStreams.start();
```

Under the hood, the “state store”

- Store / Query data
- Automatic with count / aggregations / windowing
- Use rocksDB by default
- Global state spread across stream applications
- Also resistant to failure
 - State can be rebuilt from a specific topic
 - Compacted so small footprint

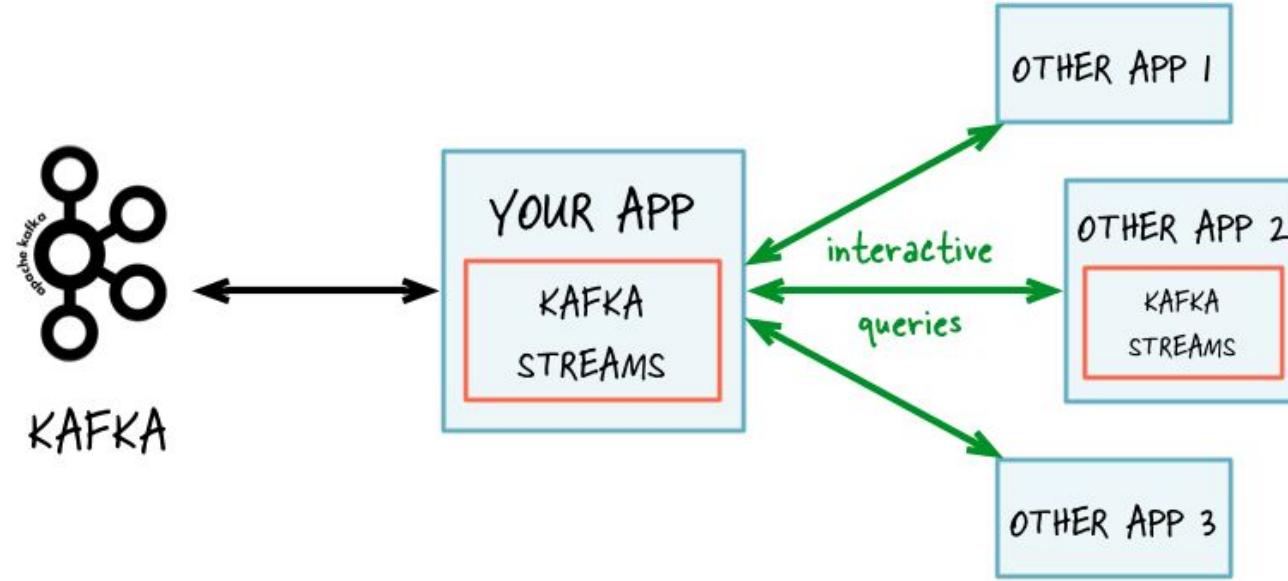


Bonus : Interactive queries

1 Capture business events in Kafka

2 Process the events with Kafka Streams

3 With interactive queries, other apps can directly query the latest results



Bonus : KSQL

- Manipulate streaming data with SQL
- Experimental

```
CREATE STREAM pageview (viewtime bigint, pageid varchar, userid varchar)
    WITH (value_format = 'json', kafka_topic='pageview_topic_json');

CREATE TABLE users (registertime bigint, userid varchar, regionid varchar, gender varchar)
    WITH (value_format = 'json', kafka_topic='user_topic_json');

-- Find the pageviews from reagion with id ending in _8 and _9 from the female pageview
CREATE STREAM enrichedpv_female_r8 AS
    SELECT *
    FROM enrichedpv_female
    WHERE regionid LIKE '%_8' OR regionid LIKE '%_9';
```

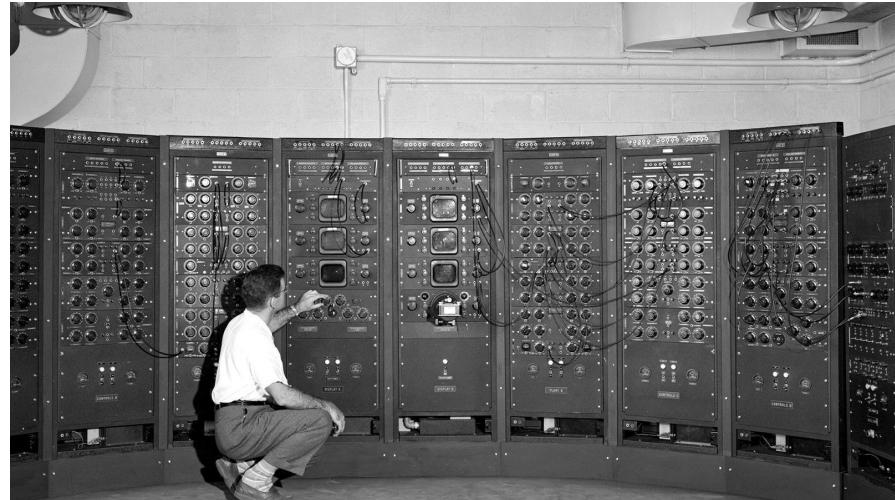


Kafka from the trenches



Basic setup

- Java 8
- RAM
- Unix systems
- Separate drives for
 - Operating System
 - Kafka logs
 - Zookeeper (or on another cluster)



Monitoring

- **Not trivial**
 - Cluster status
 - Topic health / lag
- **JMX based**
 - Replication
 - Broker status
 - Consumer / producer throughput
- **Monitoring for Kafka and Zookeeper**



Tools

- **Monitoring:**

- Confluent Control Center (\$)
- SPM (\$)
- Burrow

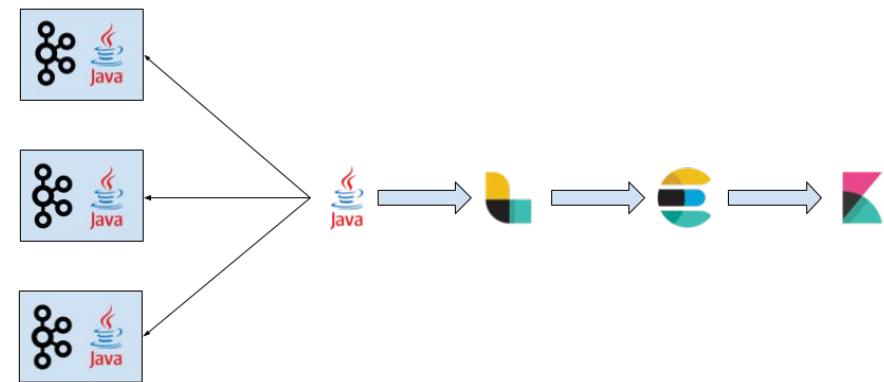
- **Management**

- Kafka Manager
- Kafkat



Monitoring : Build your own

- Java monitoring agent on each broker
 - JMX collect and simple transform
 - Expose through HTTP
- Logstash for data ingestion
- Query through Elasticsearch
- Visualize in Kibana





Conclusion

PPON

10 commandments

1. Be kind with Zookeeper
2. Exactly once you will produce
3. Big messages you shall not push
4. Tune producer & consumer settings you will
5. Size your topics
6. Compress if you need to
7. Compaction you will consider
8. Monitor and do not overreact
9. You shall stream
10. Not a silver bullet you shall consider it



Ressources

- **Examples**

- <https://github.com/lmFlog/kafka-playground>
- <https://github.com/confluentinc/examples>

- **Documentation**

- <http://docs.confluent.io/>
- <https://kafka.apache.org/>

- **Blogs / articles**

- <https://www.confluent.io/blog>
- <https://engineering.linkedin.com/distributed-systems/log-what-every-software-engineer-should-know-about-real-time-datas-unifying>
- <https://cwiki.apache.org/confluence/display/KAFKA/KIP-98---Exactly+Once+Delivery+and+Transactional+Messaging>
- <https://zcox.wordpress.com/2017/01/16/updating-materialized-views-and-caches-using-kafka/>
- <https://technology.amis.nl/2017/02/12/apache-kafka-streams-running-top-n-grouped-by-dimension-from-and-to-kafka-topic/>



PARIS
BORDEAUX
NANTES
LYON
MARRAKECH
WASHINGTON DC
NEW-YORK
RICHMOND
MELBOURNE

contact@ippon.fr
www.ippon.fr - www.ippon-hosting.com - www.ippon-digital.fr
@ippontech

-
01 46 12 48 48